



Chapter 10: Big Data

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Motivation

- Very large volumes of data being collected
 - Driven by growth of web, social media, and more recently internet-of-things
 - Web logs were an early source of data
 - Analytics on web logs has great value for advertisements, web site structuring, what posts to show to a user, etc
- Big Data: differentiated from data handled by earlier generation databases
 - **Volume**: much larger amounts of data stored
 - **Velocity**: much higher rates of insertions
 - **Variety**: many types of data, beyond relational data



Querying Big Data

- Transaction processing systems that need very high scalability
 - Many applications willing to sacrifice ACID (Atomicity, Consistency, Integrity, Durability) properties and other database features, if they can get very high scalability
- Query processing systems that
 - Need very high scalability, and
 - Need to support non-relation data



Big Data Storage Systems

- Distributed file systems
- Sharding across multiple databases
- Key-value storage systems
- Parallel and distributed databases




Distributed File Systems

- A distributed file system stores data across a large collection of machines, but provides single file-system view
- Highly scalable distributed file system for large data-intensive applications.
 - E.g., 10K nodes, 100 million files, 10 PB
- Provides redundant storage of massive amounts of data on cheap and unreliable computers
 - Files are replicated to handle hardware failure
 - Detect failures and recovers from them
- Examples:
 - Google File System (GFS)
 - Hadoop File System (HDFS)



Hadoop File System Architecture

- Single Namespace for entire cluster
- Files are broken up into blocks
 - Typically 64 MB block size
 - Each block replicated on multiple DataNodes
- Client
 - Finds location of blocks from NameNode
 - Accesses data directly from DataNode





Hadoop Distributed File System (HDFS)

- **NameNode**
 - Maps a filename to list of Block IDs
 - Maps each Block ID to DataNodes containing a replica of the block
- **DataNode**: Maps a Block ID to a physical location on disk
- Data Coherency
 - Write-once-read-many access model
 - Client can only append to existing files
- Distributed file systems good for millions of large files
 - But have very high overheads and poor performance with billions of smaller tuples



Sharding

- **Sharding**: partition data across multiple databases
- Partitioning usually done on some ***partitioning attributes*** (also known as ***partitioning keys*** or ***shard keys*** e.g. user ID
 - E.g., records with key values from 1 to 100,000 on database 1, records with key values from 100,001 to 200,000 on database 2, etc.
- Application must track which records are on which database and send queries/updates to that database
- Positives: scales well, easy to implement
- Drawbacks:
 - Not transparent: application has to deal with routing of queries, queries that span multiple databases
 - When a database is overloaded, moving part of its load out is not easy
 - Chance of failure more with more databases
 - need to keep replicas to ensure availability, which is more work for application



Parallel Databases and Data Stores

- Supporting scalable data access
 - Approach 1: memcache or other caching mechanisms at application servers, to reduce database access
 - Limited in scalability
 - Approach 2: Partition (“shard”) data across multiple separate database servers
 - Approach 3: Use existing parallel databases
 - Historically: parallel databases that can scale to large number of machines were designed for decision support not OLTP
 - Approach 4: Massively Parallel Key-Value Data Store
 - Partitioning, high availability etc. completely transparent to application
- Sharding systems and key-value stores don’t support many relational features, such as joins, integrity constraints, etc., across partitions.



Key Value Storage Systems

- Key-value storage systems store large numbers (billions or even more) of small (KB-MB) sized records
- Records are **partitioned** across multiple machines and
- Queries are routed by the system to appropriate machine
- Records are also **replicated** across multiple machines, to ensure availability even if a machine fails
 - Key-value stores ensure that updates are applied to all replicas, to ensure that their values are **consistent**



Key Value Storage Systems

- Key-value stores may store
 - **uninterpreted bytes**, with an associated key
 - E.g., Amazon S3, Amazon Dynamo
 - **Wide-table** (can have arbitrarily many attribute names) with associated key
 - Google BigTable, Apache Cassandra, Apache Hbase, Amazon DynamoDB
 - Allows some operations (e.g., filtering) to execute on storage node
 - JSON
 - MongoDB, CouchDB (document model)
- **Document stores** store semi-structured data, typically JSON
- Some key-value stores support multiple versions of data, with timestamps/version numbers



Data Representation

- An example of a JSON object is:

```
{  
  "ID": "22222",  
  "name": {  
    "firstname": "Albert",  
    "lastname": "Einstein"  
  },  
  "deptname": "Physics",  
  "children": [  
    { "firstname": "Hans", "lastname": "Einstein" },  
    { "firstname": "Eduard", "lastname": "Einstein" }  
  ]  
}
```



Key Value Storage Systems

- Key-value stores support
 - **put**(key, value): used to store values with an associated key,
 - **get**(key): which retrieves the stored value associated with the specified key
 - **delete**(key) -- Remove the key and its associated value
- Some systems also support **range queries** on key values
- Document stores also support queries on non-key attributes
 - See book for MongoDB queries
- Key value stores are not full database systems
 - Have no/limited support for transactional updates
 - Applications must manage query processing on their own
- Not supporting above features makes it easier to build scalable data storage systems
 - Also called **NoSQL** systems



Replication and Consistency

- **Availability** (system can run even if parts have failed) is essential for parallel/distributed databases
 - Via replication, so even if a node has failed, another copy is available
- **Consistency** is important for replicated data
 - All live replicas have same value, and each read sees latest version
 - Often implemented using majority protocols



Replication and Consistency

- Very large systems will partition at some point
 - Choose one of consistency or availability
- Traditional database choose consistency
- Most Web applications choose availability
 - Except for specific parts such as order processing



The MapReduce Paradigm

- Platform for reliable, scalable parallel computing
- Abstracts issues of distributed and parallel environment from programmer
 - Programmer provides core logic (via `map()` and `reduce()` functions)
 - System takes care of parallelization of computation, coordination, etc.
- Paradigm dates back many decades
 - Google Map Reduce, Hadoop, ..
- Data storage/access typically done using distributed file systems or key-value stores



MapReduce: Word Count Example

- Consider the problem of counting the number of occurrences of each word in a large collection of documents
- How would you do it in parallel?
- Solution:
 - Divide documents among workers
 - Each worker parses document to find all words, map function outputs (word, count) pairs
 - Partition (word, count) pairs across workers based on word
 - For each word at a worker, reduce function locally add up counts
- Given input: “One a penny, two a penny, hot cross buns.”
 - Records output by the map() function would be
 - (“One”, 1), (“a”, 1), (“penny”, 1), (“two”, 1), (“a”, 1), (“penny”, 1), (“hot”, 1), (“cross”, 1), (“buns”, 1).
 - Records output by reduce function would be
 - (“One”, 1), (“a”, 2), (“penny”, 2), (“two”, 1), (“hot”, 1), (“cross”, 1), (“buns”, 1)



Pseudo-code of Word Count

map(String record):

 for each word in record

 emit(word, 1);

// First attribute of emit above is called **reduce key**

// In effect, group by is performed on reduce key to create a

// list of values (all 1's in above code). This requires **shuffle step**

// across machines.

// The reduce function is called on list of values in each group

reduce(String key, List value_list):

 String word = key

 int count = 0;

 for each value in value_list:

 count = count + value

 Output(word, count);



MapReduce Programming Model

- Inspired from map and reduce operations commonly used in functional programming languages.
- Input: a set of key/value pairs
- User supplies two functions:
 - **map**(k,v) $\rightarrow$ list(k1,v1)
 - **reduce**(k1, list(v1)) $\rightarrow$ v2
- (k1,v1) is an intermediate key/value pair
- Output is the set of (k1,v2) pairs
- For our example, assume that system
 - Breaks up files into lines, and
 - Calls map function with value of each line



MapReduce Example 2: Log Processing

- Given log file in following format:

...

2013/02/21 10:31:22.00EST [/slide-dir/11.ppt](#)

2013/02/21 10:43:12.00EST [/slide-dir/12.ppt](#)

2013/02/22 18:26:45.00EST [/slide-dir/13.ppt](#)

2013/02/22 20:53:29.00EST [/slide-dir/12.ppt](#)

...

- Goal: find how many times each of the files in the slide-dir directory was accessed between 2013/01/01 and 2013/01/31.
- Options:
 - Sequential program too slow on massive datasets
 - Load into database expensive, direct operation on log files cheaper
 - Custom built parallel program for this task possible, but very laborious
 - Map-reduce paradigm



MapReduce: File Access Count Example


```
map(String key, String record) {
    String attribute[3];
    .... break up record into tokens (based on space character), and store the
        tokens in array attributes
    String date = attribute[0];
    String time = attribute[1];
    String filename = attribute[2];
    if (date between 2013/01/01 and 2013/01/31
        and filename starts with "/slide-dir/")
        emit(filename, 1).
}

reduce(String key, List recordlist) {
    String filename = key;
    int count = 0;
    For each record in recordlist
        count = count + 1.
    output(filename, count)
}
```




Schematic Flow of Keys and Values

- Flow of keys and values in a map reduce task





Parallel Processing of MapReduce Job





Hadoop MapReduce

- Hadoop is a widely used open source implementation of Map Reduce written in Java
 - Map and reduce functions can be written in several different languages, we use Java.
- Input and output to map reduce systems such as Hadoop must be done in parallel
 - Hadoop uses Hadoop File System (HDFS),
 - Input files can be in several formats
 - Text/CSV
 - compressed representation such as Avro, ORC and Parquet
 - Hadoop also supports key-value stores such as Hbase, Cassandra, MongoDB, etc.



Data Structures in Hadoop

- Generic Mapper and Reducer interfaces both take four type arguments, that specify the types of the
 - input key, input value, output key and output value
- Map class in next slide implements the Mapper interface
 - Map input key is of type LongWritable, i.e. a long integer
 - Map input value which is (all or part of) a document, is of type Text.
 - Map output key is of type Text, since the key is a word,
 - Map output value is of type IntWritable, which is an integer value.



Hadoop Code in Java: Map Function

```
public static class Map extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException
    {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            context.write(word, one);
        }
    }
}
```



Hadoop Code in Java: Reduce Function

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {  
    public void reduce(Text key, Iterable<IntWritable> values,  
        Context context) throws IOException, InterruptedException  
    {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        context.write(key, new IntWritable(sum));  
    }  
}
```



Hadoop Job Parameters

- The classes that contain the map and reduce functions for the job
 - Set by methods `setMapperClass()` and `setReducerClass()`
- The types of the job's output key and values
 - Set by methods `setOutputKeyClass()` and `setOutputValueClass()`
- The input format of the job
 - Set by method `job.setInputFormatClass()`
 - Default input format in Hadoop is the `TextInputFormat`,
 - Map key whose value is a byte offset into the file, and
 - Map value is the contents of one line of the file
- The directories where the input files are stored, and where the output files must be created
 - Set by `addInputPath()` and `addOutputPath()`
- And many more parameters



Hadoop Code in Java: Overall Program

```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        Job job = new Job(conf, "wordcount");  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        job.setMapperClass(Map.class);  
        job.setReducerClass(Reduce.class);  
        job.setInputFormatClass(TextInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        job.waitForCompletion(true);  
    }  
}
```



Map Reduce vs. Databases

- Map Reduce widely used for parallel processing
 - Google, Yahoo, and 100's of other companies
 - Example uses: compute PageRank, build keyword indices, do data analysis of web click logs,
 - Allows procedural code in map and reduce functions
 - Allows data of any type
- Many real-world uses of MapReduce cannot be expressed in SQL
- But many computations are much easier to express in SQL
 - Map Reduce is cumbersome for writing simple queries



Map Reduce vs. Databases (Cont.)

- Relational operations (select, project, join, aggregation, etc.) can be expressed using Map Reduce
- SQL queries can be translated into Map Reduce infrastructure for execution
 - Apache Hive SQL, Apache Pig Latin, Microsoft SCOPE
- Current generation execution engines support not only Map Reduce, but also other algebraic operations such as joins, aggregation, etc. natively.



STREAMING DATA



Streaming Data and Applications

- **Streaming data** refers to data that arrives in a continuous fashion
 - Contrast to **data-at-rest**
- Applications include:
 - Stock market: stream of trades
 - e-commerce site: purchases, searches
 - Sensors: sensor readings
 - Internet of things
 - Network monitoring data
 - Social media: tweets and posts can be viewed as a stream
- Queries on streams can be very useful
 - Monitoring, alerts, automated triggering of actions



Querying Streaming Data

Approaches to querying streams:

- **Windowing**: Break up stream into windows, and queries are run on windows
 - Stream query languages support window operations
 - Windows may be based on time or tuples
 - Must figure out when all tuples in a window have been seen
 - Easy if stream totally ordered by timestamp
 - **Punctuations** specify that all future tuples have timestamp greater than some value
- **Continuous Queries**: Queries written e.g. in SQL, output partial results based on stream seen so far; query results updated continuously
 - Have some applications, but can lead to flood of updates



Querying Streaming Data (Cont.)

Approaches to querying streams (Cont.):

- **Algebraic operators on streams:**
 - Each operator consumes tuples from a stream and outputs tuples
 - Operators can be written e.g., in an imperative language
 - Operator may maintain state
- **Pattern matching:**
 - Queries specify patterns, system detects occurrences of patterns and triggers actions
 - E.g., Microsoft StreamInsight, Flink CEP, Oracle Event Processing



Stream Processing Architectures

- Many stream processing systems are purely in-memory, and do not persist data
- **Lambda architecture**: split stream into two, one output goes to stream processing system and the other to a database for storage
 - Easy to implement and widely used
 - But often leads to duplication of querying effort, once on streaming system and once in database



Stream Extensions to SQL

- Streaming systems often support more window types
 - **Tumbling window**
 - E.g., hourly windows, windows don't overlap
 - **Hopping window**
 - E.g., hourly window computed every 20 minutes
 - **Sliding window**
 - Window of specified size (based on timestamp interval or number of tuples) around each incoming tuple
 - **Session window**
 - Groups tuples based on user sessions



Window Syntax in SQL

- Windowing syntax varies widely by system
- E.g., in Azure Stream Analytics SQL:


```
select item, System.Timestamp as window end, sum(amount)  
from order timestamp by datetime  
group by itemid, tumblingwindow(hour, 1)
```

- Aggregates are applied on windows
- Result of windowing operation on a stream is a relation
- Many systems support stream-relation joins
- Stream-stream joins often require join conditions to specify bound on timestamp gap between matching tuples
 - E.g., tuples must be at most 30 minutes apart in timestamp




Algebraic Operations on Streams

- Tuples in streams need to be routed to operators
- Routing of streams using DAG and publish-subscribe representations
 - Used in Apache Storm and Apache Kafka respective



(a) DAG representation of streaming data flow



(b) Publish-subscribe representation of streaming data flow



GRAPH DATABASES



Graph Data Model

- Graphs are a very general data model
- ER model of an enterprise can be viewed as a graph
 - Every entity is a node
 - Every binary relationship is an edge
 - Ternary and higher degree relationships can be modelled as binary relationships



Graph Data Model (Cont.)

- Graphs can be modelled as relations
 - *node(ID, label, node_data)*
 - *edge(fromID, toID, label, edge_data)*
- Above representation too simplistic
- Graph databases like Neo4J can provide a **graph view of relational schema**
- Query languages for graph databases make it
 - easy to express queries requiring edge traversal
 - allow efficient algorithms to be used for evaluation



Graph Data Model (Cont.)

- Suppose
 - Relations *instructor* and *student* are nodes, and
 - Relation *advisor* represents edges between instructors and student
- Query in Neo4J:
match (*i:instructor*)<-[:*advisor*]-(*s:student*)
where *i.dept name*= 'Comp. Sci.'
return *i.ID* **as** ID, *i.name* **as** name, **collect**(*s.name*) **as** advisees
- **match** clause matches nodes and edges in graphs



Parallel Graph Processing

- Very large graphs (billions of nodes, trillions of edges)
 - Web graph: web pages are nodes, hyper links are edges
 - Social network graph: people are nodes, friend/follow links are edges
- Two popular approaches for parallel processing on such graphs
 - Map-reduce and algebraic frameworks
 - **Bulk synchronous processing (BSP)** framework
- Multiple iterations are required for any computations on graphs
 - Map-reduce/algebraic frameworks often have high overheads per iteration
 - BSP frameworks have much lower per-iteration overheads
- Google's Pregel system popularized the BSP framework
- Apache Giraph is an open-source version of Pregel
- Apache Spark's GraphX component provides a Pregel-like API



Bulk Synchronous Processing

- Each vertex (node) of a graph has data (state) associated with it
 - Vertices are partitioned across multiple machines, and state of node kept in-memory
- Analogous to map() and reduce() functions, programmers provide methods to be executed for each node
 - Node method can send messages to or receive messages from neighbouring nodes
- Computation consists of multiple iterations, or supersteps



End of Chapter 10